

# Out of the Box Replication In Postgres 9.4



**Denish Patel**  
Lead Database Architect

## Who am I ?

- Database Architect with **OmniTI** for last 8+ years
- Expertise in PostgreSQL , Oracle, MySQL, NoSQL
- Contact : [denish@omniti.com](mailto:denish@omniti.com)
- Twitter: @DenishPatel
- Blog: <http://www.pateldenish.com>
- Providing Solutions for business problems to deliver
  - Scalability
  - Reliability
  - High Availability
  - Security

**We are hiring!!**  
Apply @ [l42.org/lg](http://l42.org/lg)

# Agenda



- What is WAL?
- Postgres Replication History
- How you setup replication now?
- What's missing ?
- Why replication slots?
- Demo

I'm **NOT** going to discuss ...

- SLONY or other 3<sup>rd</sup> party replication tools
- 3<sup>rd</sup> party replication management tools

## WAL - Write Ahead Log

- Roll-forward recovery aka REDO
- flushed to disk to guarantee commit durability
- sequential writes
- lower cost than flushing page cache
- Allows us to do cool things
  - Crash recovery
  - Binary backups
  - Point-In Time Recovery
  - Replication

## WAL Internals (Basic)

- Automatically enabled; no action required
- Make sure to have enough space on server
- Stored under pg\_xlog directory
- Normally, 16MB in size
  - ✓ --with-wal-segsize config option at build
- Each segment is divided into pages (8 kB page)
  - ✓ --with-wal-blocksize config option at build
- Segment name starts with..  
00000000100000000000000000000000

## What is logged?

```
select * from pg_settings where name='wal_level';
```

name	<b>wal_level</b>
setting	hot_standby
unit	
category	Write-Ahead Log / Settings
short_desc	Set the level of information written to the WAL.
extra_desc	
context	postmaster
vartype	enum
source	configuration file
min_val	
max_val	
enumvals	{ <b>minimal</b> , <b>archive</b> , hot_standby, <b>logical</b> }
boot_val	minimal
reset_val	hot_standby
sourcefile	/var/lib/pgsql/9.4/data/postgresql.auto.conf
sourceline	4

## What's **NOT** logged?

**Almost everything is replicated, but...**

- unlogged tables (As name suggests)
- temporary tables
- hash indexes? (generally don't use?)

# Postgres Replication History

Postgres 7.0: WAL

Postgres 8.0: PITR (Point-In-Time-Recovery)

Postgres 8.2: pg\_standby

Postgres 9.0: Hot\_standby, Streaming replication

Postgres 9.1: pg\_basebackup, Synchronous replication

Postgres 9.2: Cascading Replication

Postgres 9.3: Standby can switch timeline to follow new master

Postgres 9.4: Replication Slots , Logical decoding



# Basic Steps to Setting up Replication

initdb

# Postgresql.conf

- `max_wal_senders=10`
- `wal_level=hot_standby`
- `hot_standby=on` (on standby)

## pg\_hba.conf

#TYPE	DATABASE	USER	ADDRESS	METHOD
host	replication	replication	10.0.0.1/32	md5

## Restart database ...



```
pg_ctl restart
```

## Create replication user



```
CREATE ROLE replication WITH LOGIN  
REPLICATION;
```

```
\password replication
```

# Take Backup

- Take file system level backups
- What tools you are using for backups?

## Recovery.conf



```
primary_conninfo = 'host=primaryhost user=replication  
password=replication'
```

```
standby_mode = on
```

## Startup standby db

pg\_ctl start



Not done yet!

Have you configured archiving?

Don't forget about this setting?

wal\_keep\_segments

## Setup archiving . . .

- postgresql.conf  
archive\_mode = on  
archive\_command = 'cp %p /some/where/%f'

## Setup restore command

- recovery.conf

```
restore_command = 'cp /some/where/%f %p'
```

# Archiving Options



- local or remote copy
- Scp
- Rsync
- NFS
- pg\_archivecleanup

## What if you have more standby machines?

```
archive_command = 'rsync %p standby1::pg/%f && rsync  
%p standby2::pg/%f'
```

```
archive_command = 'echo standby1 standby2 ...  
| xargs -d" " -I{} -n1 -Po -r rsync %p {}::pg/%f'
```

# Replication management tools?

- OmniPITR
- WAL-E
- Repmgr
- Pgbarman
- Skytools
- A lot of Custom scripts

# What about fsync?

- **fsync** capabilities
  - cp: no
  - dd: GNU coreutils
  - SSH: OpenSSH 6.5 sftp-server (Jan 2014)
  - rsync: patch or wrapper
  - NFS: Supported



## Postgres 9.4 ; Enter Replication Slots

# Postgresql.conf



```
max_replication_slots = 8
```

## Create Slot

```
SELECT * FROM  
pg_create_physical_replication_slot('name');
```

## Primary knows the status of standbys

```
select * from pg_replication_slots ;
```

```
slot_name | plugin | slot_type | datoid | database | active | xmin | catalog_xmin | restart_lsn
```

```
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----  
standby1 |      | physical |      |      | t      |      |      | o/21000058  
standby2 |      | physical |      |      | t      |      |      | o/21000058  
standby3 |      | physical |      |      | t      |      |      | o/21000058
```

## Recovery.conf



```
primary_slot_name = 'name'
```

## Replication slots benefits

- Keep necessary WAL files
- Each standby can have different WAL apply status
- Single access control setup
- fsync on receiving side

## pg\_basebackup

```
pg_basebackup \  
-h primaryhost \  
-U replication \  
-D $PGDATA \  
-X stream \  
-P -v -R
```

## How about archiving?

- Meet `pg_receivexlog`
  - (Available since Postgres 9.1!)

`pg_receivexlog \`

`-D archivedir \`

`--slot archiving_slot \`

`-h primaryhost -U replication`

- `-- synchronous` option in 9.5



## Postgres 9.4 – Out of the Box Replication

Are you ready to setup replication without any external tools?

1. `pg_basebackup`
2. streaming with Replication Slots
3. `pg_receivexlog`

## Tutorial – Let's bring up VM!

- Google Drive: <https://drive.google.com/open?id=oBxnXwkT5PRBdeVByQVYySkhIemc>
- Login: **pgtraining/pgcon**
- pgtraining user has *sudo* access
- You can access *terminal* on the desktop
- **Internet** should be working within VM
- **Copy/paste** should work between VM and Host
- Take **snapshot** along the process so you can rollback easily

# Installing Postgres

- Remove Postgres 8.4 version

```
sudo yum erase postgresql.*
```

- Setup yum repo for your desired version

```
sudo yum install http://yum.postgresql.org/9.4/redhat/rhel-6-x86_64/pgdg-redhat94-9.4-1.noarch.rpm
```

- Install PostgreSQL 9.4 & Contrib modules

```
sudo yum install postgresql94-server postgresql94-contrib
```

- Create postgres cluster & initial automatic startup

```
sudo service postgresql-9.4 initdb
```

```
sudo chkconfig postgresql-9.4 on
```

```
sudo service postgresql-9.4 start
```

# Create Role and Database



- Become postgres system user

`sudo su - postgres`

- Log into database using psql (\? to see all available commands)
- Create a role & database for yourself

`CREATE ROLE pgtraining WITH LOGIN SUPERUSER;`

`CREATE DATABASE pgtraining;`

- You can login as pgtraining user (`psql -U pgtraining -d pgtraining`)
- Create replication role for later

`CREATE ROLE replication WITH LOGIN REPLICATION;`

- Set password ("replication")

`\password replication`

## Configure pg\_hba.conf

- <http://www.postgresql.org/docs/9.4/static/auth-pg-hba-conf.html>

- Find location of hba\_file

```
postgres=# show hba_file;
```

- Add following entry for replication user
- Try to avoid **trust** authentication
- [pgtraining@localhost ~]\$ sudo vi /var/lib/pgsql/9.4/data/pg\_hba.conf

```
host    replication    replication    127.0.0.1/32    md5
```

## Prepare Primary DB server

```
alter system set wal_level = hot_standby;  
alter system set archive_mode=on;  
alter system set max_replication_slots=8;  
alter system set archive_timeout = 60;  
alter system set max_wal_senders = 8;  
alter system set wal_keep_segments=100;  
alter system set logging_collector=on;
```

## Restart Primary DB server

- Restart database

```
sudo service postgresql-9.4 restart
```

- Verify settings

```
psql=# show max_wal_senders;
```

## Create replication slot

```
SELECT * FROM pg_create_physical_replication_slot('standby1');
```



## Let's take backup

```
sudo su - postgres
```

```
pg_basebackup -h 127.0.0.1 -U replication -D /var/  
lib/pgsql/9.4/slave -R -Xs -P -v
```

## Prepare standby db server ...

```
cd /var/lib/pgsql/9.4/slave
```

- Edit **Standby** postgresql.conf file

```
port = 5433
```

```
hot_standby = on
```

- Edit **Standby** recovery.conf file

```
standby_mode = 'on'
```

```
primary_conninfo = 'user=replication password=replication  
host=127.0.0.1 port=5432'
```

```
primary_slot_name='standby1'
```

```
trigger_file = '/var/lib/pgsql/9.4/slave/finish.recovery'
```

```
recovery_target_timeline='latest'
```

## Configure Standby with init

- Copy existing init file

```
sudo cp /etc/init.d/postgresql-9.4 /etc/init.d/postgresql-9.4-5433
```

- Edit config file to change (as root):

```
PGDATA=/var/lib/pgsql/9.4/slave
```

```
PGLOG=/var/lib/pgsql/9.4/pgstartup-5433.log
```

```
PGUPLOG=/var/lib/pgsql/$PGMAJORVERSION/  
pgupgrade-5433.log
```

- Register service, start up slave

```
sudo chkconfig postgresql-9.4-5433 on
```

```
sudo service postgresql-9.4-5433 start
```

## Status: pg\_stat\_replication



pgtraining=# \x

**pgtraining=# select \* from pg\_stat\_replication;**

```
-[ RECORD 1 ]-----+-----  
pid          | 3260  
usesysid     | 24576  
username     | replication  
application_name | walreceiver  
client_addr  | 127.0.0.1  
client_hostname |  
client_port  | 53206  
backend_start | 2015-06-08 14:47:50.057326-04  
backend_xmin  |  
state        | streaming  
sent_location | 0/240000B8  
write_location | 0/240000B8  
flush_location | 0/240000B8  
replay_location | 0/240000B8  
sync_priority | 0  
sync_state   | async
```

## Status : pg\_replication\_slots



```
pgtraining=# select * from pg_replication_slots;
```

```
-[ RECORD 1 ]+-----
```

```
slot_name    | standby1
```

```
plugin       |
```

```
slot_type    | physical
```

```
datoid       |
```

```
database     |
```

```
active       | t
```

```
xmin         |
```

```
catalog_xmin |
```

```
restart_lsn  | 0/270000EC
```

## What about archiving?

- Create slot for archiving:

```
SELECT * FROM  
pg_create_physical_replication_slot('archiver1');
```

- Create archive directory

```
mkdir /var/lib/pgsql/9.4/archive
```

- Start archiving process in background

```
/usr/pgsql-9.4/bin/pg_receivexlog -h 127.0.0.1 -p 5432 -U  
replication -S 'archiver1' -n -v -D /var/lib/pgsql/9.4/  
archive
```

- Put under init.d for continuous run
- Switch xlog : primary\_db\_sever# `select pg_switch_xlog();`

# Monitoring

- Monitor Disk space
- Monitor slave lag

```
select pg_xlog_location_diff(sent_location, write_location) AS  
       byte_lag
```

```
from pg_stat_replication
```

```
where application_name='pg_receivexlog';
```

- Monitor WAL archive process

```
select pg_xlog_location_diff(sent_location, write_location) AS  
       byte_lag
```

```
from pg_stat_replication
```

```
where application_name='walreceiver';
```

# Monitoring

- Number of pg\_xlogs on master
- Monitor pg\_recievevexlog process
- Make sure archive location has new files
- pg\_basebackup log files for successful backup... look for “pg\_basebackup: base backup completed”



## Thanks to....

- OmniTI for travel sponsorships
- Pgcon conference committee
- Peter Eisentraut
- You!!

# Questions?

[denish@omniti.com](mailto:denish@omniti.com)

Twitter: DenishPatel